
Evaluating Self-Supervised Pretraining Without Using Labels

Colorado Reed^{†*} Sean Metzger^{‡*} Aravind Srinivas[†] Trevor Darrell[†] Kurt Keutzer[‡]

[†]BAIR, Department of Computer Science, UC Berkeley

[‡]Graduate Group in Bioengineering (Berkeley/UCSF),
Weill Neurosciences Institute & UCSF Neurological Surgery

Abstract

A common practice in unsupervised representation learning is to use labeled data to evaluate the learned representations — oftentimes using the labels from the “unlabeled” training dataset. This supervised evaluation is then used to guide the training process, e.g. to select augmentation policies. However, supervised evaluations may not be possible when labeled data is difficult to obtain (such as medical imaging) or ambiguous to label (such as fashion categorization). This raises the question: is it possible to evaluate unsupervised models without using labeled data? Furthermore, is it possible to use this evaluation to make decisions about the training process, such as which augmentation policies to use? In this work, we show that the simple self-supervised evaluation task of image rotation prediction is highly correlated with the supervised performance of standard visual recognition tasks and datasets (rank correlation > 0.94). We establish this correlation across hundreds of augmentation policies and training schedules and show how this evaluation criteria can be used to automatically select augmentation policies without using labels. Despite not using any labeled data, these policies perform comparably with policies that were determined using supervised downstream tasks. Importantly, this work explores the idea of using unsupervised evaluation criteria to help both researchers and practitioners make decisions when training without labeled data.

1 Introduction

Self-supervised learning, a type of unsupervised learning that creates target objectives without human annotation, has led to a dramatic increase in the ability to capture salient feature representations from unlabeled visual data. So much so, that in an increasing number of cases, these feature representations can outperform the representations learned from the same data with labels [1, 2, 3]. At the center of these recent advances is a form of *instance contrastive learning* where a single image is augmented using two separate data augmentations, and then a network is trained to distinguish which augmented images originated from the same image when contrasted with other augmented images, see [2, 3, 4, 5].

As illustrated in Figure 1, recent works such as [2, 4] have used extensive supervised evaluations to determine which augmentation policies to use for training. The best policies obtain a *sweet spot*, where the image transformations make it difficult for the contrastive task to determine the corresponding image pairs, while still retaining salient image features; finding this sweet spot can be the difference between state-of-the-art performance or poor performance for various tasks. However, it is often difficult or impossible to obtain accurately labeled data in privacy sensitive fields (e.g. medical imaging [6]) or applications with highly ambiguous label definitions (e.g. fashion categorization [7]). This leads to the open question: *How can we evaluate self-supervised models, especially to select augmentation policies, when labeled data is not available?* We address this question as follows:

*Equal contribution. Please send correspondence to cjrd@cs.berkeley.edu.

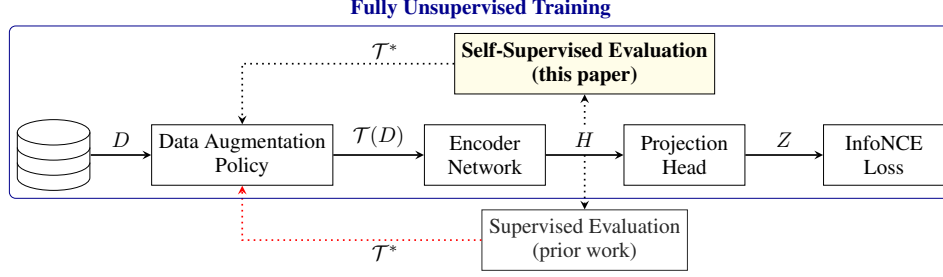


Figure 1: The blue box highlights our fully unsupervised training pipeline for instance contrastive representation learning: data D are augmented with policy \mathcal{T} , then transformed into representations, H , which are fed into a projection head yielding features, Z , that are used to determine the InfoNCE loss (Eq. 1). As shown by the red arrow, prior work uses supervised evaluations of the representations, F to inform the training process, e.g. to update the augmentation policy $\mathcal{T} \rightarrow \mathcal{T}^*$. In this paper, we show that a simple self-supervised evaluation criteria can be used in lieu of the supervised evaluation, and that this can lead to automated and efficient selection of augmentations.

- We show that an image rotation prediction evaluation task is highly correlated with the downstream supervised performance (rank correlation $\rho > 0.94$) on six standard recognition datasets and tasks across hundreds of learned representations, spanning three types of common evaluation techniques: linear separability performance, semi-supervised performance, and transfer learning performance.
- Using self-supervised evaluation, we adapt two automatic data augmentation algorithms for instance contrastive learning. These algorithms efficiently discover augmentation policies without the use of any labels that match or outperform policies obtained using supervised feedback.

We conclude that image rotation prediction is a strong, unsupervised evaluation criteria for evaluating and selecting data augmentations for instance contrastive learning. We hope this work will encourage future research to leverage unsupervised evaluation criteria, as it opens up the wealth of supervised techniques that have been developed for making decisions when building and training models.

2 Related work

In this paper, we study the evaluation techniques employed for self-supervised representations, particularly through the lens of learning data augmentation policies. We discuss these topics next.

Self-supervised representation learning: The general goal of representation learning is to pretrain a network and then either fine-tune it for a particular task or transfer it to a related model, e.g. see [1, 3, 8, 9, 10, 11, 12, 13, 14]. Recently, [2] and [4] demonstrated substantial improvements by using similar forms of instance contrastive learning whereby one image is augmented using two separate data augmentations and then a network is trained to identify this positive pair contrasted with a large set of distractor images. The most commonly used contrastive loss function is the InfoNCE loss, where given two images originating from the same image i and K_d distractor images we have:

$$\mathcal{L}_{NCE} = -\mathbb{E} \left[\log \frac{\exp(\text{sim}(\mathbf{z}_{1,i}, \mathbf{z}_{2,i}))}{\sum_{j=1}^{K_d} \exp(\text{sim}(\mathbf{z}_{1,i}, \mathbf{z}_{2,j}))} \right] \quad (1)$$

where $\mathbf{z}_{1,i}, \mathbf{z}_{2,i}$ are the two different image representations from image i following the encoder network and projection head as shown in Figure 1, and $\text{sim}(\cdot, \cdot)$ is a similarity function such as a weighted dot product. The InfoNCE loss [15, 1] has been shown to maximize a lower bound on the mutual information $I(\mathbf{h}_1; \mathbf{h}_2)$. The SimCLR framework [2] relies on large batch sizes to contrast the image pairs, while the MoCo framework [4] maintains a large, continually-updating queue of contrasting images. Given the increased adoption, broad application, and strong performance of instance contrastive learning [1, 2, 4], we focus our work on this type of self-supervised learning.

Evaluating self-supervised models: Evaluating self-supervised models is typically done via:

- *Separability*: the network is frozen and the training data is used to train a supervised linear model; the justification is that good representations will reveal linear separability in the data [15, 16, 17, 18, 19, 20, 21].
- *Transferability*: the network is either frozen or jointly fine-tuned, with a transfer task model that is fine-tuned using a different, labeled dataset [1, 3, 13, 14]; the justification is that good representations will generalize to a number of downstream tasks.
- *Semi-supervised*: the network is either frozen or jointly fine-tuned using a fraction of labeled data with either the *separability* or *transferability* tasks mentioned above [1, 2]; the justification is that good representations will strongly benefit from a small number of labels.

While these evaluations characterize the unsupervised model in several ways, they have limited use for making training decisions because: **(i) accessing labels as a part of the training process is not possible for unlabeled datasets, and (ii) evaluating the model on a different, labeled dataset requires an integrated understanding of the relationship between the transfer task, datasets, and models. We seek a simpler, task-agnostic evaluation strategy that does not require labels or additional data/models.**

Learning data augmentation policies: Data augmentation has played a fundamental role in visual learning, and indeed, has a large body of research supporting its use, see [22]. In this work, we use a self-supervised evaluation to automatically learn an augmentation policy for instance contrastive models. To formulate our automatic data augmentation framework, we draw on several, equally competitive recent works in the supervised learning space [23, 24, 25, 26], where [23, 24, 25] use a separate search phase to determine the augmentation policy and [26] use a simplified augmentation space and sample from it via a grid search. For instance contrastive learning, we adapt a search-based automatic augmentation framework, Fast AutoAugment (FAA) [25], and a sampling-based approach, RandAugment [26]. We discuss these two algorithms in greater detail in the following section.

3 Self-supervised data augmentation selection

Our central goals are to **(i) establish a strong correlation between the evaluation performance of a self-supervised evaluation task and common recognition tasks used in evaluating self-supervised models, and (ii) develop a practical algorithm for self-supervised data augmentation selection.** The following subsections defines these goals in more detail.

3.1 Supervised and self-supervised evaluation correlation

With labeled data, augmentation policy selection can directly optimize the supervised task performance [24, 25]. With unlabeled data, we seek an evaluation criteria that is highly correlated with the supervised performance without requiring labels. In [27], the authors show that several self-supervised tasks are correlated with the supervised performance of models and can be used to drive neural architecture search. Inspired by this work, we study image rotation prediction — originally a task used for representation learning itself [17] — as a self-supervised task for evaluating self-supervised models. Specifically, we train a linear layer on top of a frozen backbone network to predict which rotation among $[0^\circ, 90^\circ, 180^\circ, 270^\circ]$ was applied to the input data. Note: we conducted several preliminary experiments with alternative self-supervised tasks and chose rotation prediction as it had the strongest performance — see Appendix D.

Why would image rotation prediction correlate with supervised performance? In [17], the authors show convolutional attention maps that indicate that rotation prediction networks focus on high-level object parts in an image, e.g. noses, mouths, tails, and heads. In other words, the network must learn distinguishable object features to determine whether the image has been rotated. We thus hypothesize that image rotation prediction with a simple, linear model can be used to evaluate whether the learned representations capture salient object-level features. We show that object-centric downstream applications strongly correlate with this evaluation, but expect that other types of applications such as texture or color recognition may require a different self-supervised evaluation task.

To establish the correlation between image rotation prediction and downstream performance, we train a MoCo network [3] using a range of augmentation strategies, training schedules, and datasets on image classification, object detection, and scene classification tasks. Following [2, 4, 28], we evaluate the representations on standard benchmark tasks used in self-supervised representation learning: **(i)**

Algorithm 1: SelfAugment

Input : $(D_{\text{train}}, K, T, B, P)$

```
1 Split  $D_{\text{train}}$  into  $K$ -folds:  $D_{\text{train}}^{(k)} = \{(D_{\mathcal{M}}^{(k)}, D_{\mathcal{A}}^{(k)})\}$  // split into  $K$  folds
2 for  $k \in \{1, \dots, K\}$  do
3    $\mathcal{T}^{*(k)} \leftarrow \emptyset, (D_{\mathcal{M}}, D_{\mathcal{A}}) \leftarrow (D_{\mathcal{M}}^{(k)}, D_{\mathcal{A}}^{(k)})$  // initialize
4   Train  $\theta_{\text{moco}}$  on  $D_{\mathcal{M}}$ 
5   Train  $\phi_{\text{rot}}$  on  $D_{\mathcal{M}}$  on top of frozen  $\theta_{\text{moco}}$ 
6   for  $t \in \{0, \dots, T-1\}$  do
7      $\mathcal{B} \leftarrow \text{BayesOptim}(\mathcal{T}, \mathcal{L}(\theta_{\text{moco}}, \phi_{\text{rot}} | \mathcal{T}(D_{\mathcal{A}})), B)$  // explore-and-exploit
8      $\mathcal{T}_t \leftarrow \text{Select top-}P \text{ policies in } \mathcal{B}$ 
9      $\mathcal{T}^{*(k)} \leftarrow \mathcal{T}^{*(k)} \cup \mathcal{T}_t$  // merge augmentation policies
10 return  $\mathcal{T}^* = \bigcup_k \mathcal{T}^{*(k)}$ 
```

training a supervised linear classifier on top of the frozen network, **(ii)** training a model for different downstream dataset and task while either freezing the network or finetuning it for the task, **(iii)** using a subset of labeled data to train a linear classifier on top of the frozen network. For (i-iii), we establish a strong correlation between the supervised performance and linear image rotation prediction accuracy.

3.2 Self-supervised data augmentation policies

We study and adapt two approaches for augmentation policy selection from the supervised domain: a sampling-based strategy, RandAugment [26], and a search-based strategy, Fast AutoAugment (FAA) [25]. Using the notation from [25], let \mathbb{O} represent the set of image transformations operations $\mathcal{O} : \mathcal{X} \rightarrow \mathcal{X}$ on input image \mathcal{X} . Following [25, 26], we define \mathbb{O} as the set (see Appendix B for more details):

- | | | | | |
|-----------|----------------|---------------|---------------|------------|
| • cutout | • autoContrast | • equalize | • rotate | • solarize |
| • color | • posterize | • contrast | • brightness | • sharpnes |
| • shear-x | • shear-y | • translate-x | • translate-y | • invert |

Each transformation \mathcal{O} has two parameters: **(i)** the magnitude λ that determines the strength of the transformation and **(ii)** the probability of applying the transformation, p . Let \mathcal{S} represent the set of augmentation sub-policies, where a sub-policy $\tau \in \mathcal{S}$ is defined as the sequential application of N_τ consecutive transformation $\{\bar{\mathcal{O}}_n^{(\tau)}(x; p_n^{(\tau)}, \lambda_n^{(\tau)}) : n = 1, \dots, N_\tau\}$ where each operation is applied to an input image sequentially with probability p . Applying sub-policy $\tau(x)$ is then a composition of transformation $\tilde{x}_{(n)} = \bar{\mathcal{O}}_n^{(\tau)}(\tilde{x}_{(n-1)})$ for $n = 1, \dots, N_\tau$, where the full sub-policy application has shorthand $\tilde{x}_{(N_\tau)} = \tau(x)$ and the first application is $\tilde{x}_{(0)} = x$. The full policy, \mathcal{T} , is a collection of $N_\mathcal{T}$ sub-policies, and $\mathcal{T}(D)$ represents the set of augmented images from D obtained by applying \mathcal{T} .

SelfRandAugment: RandAugment makes the following simplifying assumptions: **(i)** all transformations share a single, discrete magnitude, $\lambda \in [1, 30]$ **(ii)** all sub-policies apply the same number of transformations, N_τ **(iii)** all transformations are applied with uniform probability, $p = K_T^{-1}$ for the $K_T = |\mathbb{O}|$ transformations. RandAugment selects the best result from a grid search over (N_τ, λ) . To adapt this algorithm for self-supervised learning, we simply evaluate the searched (N_τ, λ) states using the linear image rotation prediction task, and refer to this as *SelfRandAugment*.

SelfAugment: We adapt the search-based FAA algorithm to the self-supervised setting; we call this adaptation *SelfAugment*. Formally, let \mathcal{D} be a distribution on the data \mathcal{X} . For model $\mathcal{M}(\cdot|\theta) : \mathcal{X}$ with parameters θ , define the supervised loss as $\mathcal{L}(\theta|D)$ on model $\mathcal{M}(\cdot|\theta)$ with data $D \sim \mathcal{D}$. For any given pair of D_{train} and D_{valid} , FAA selects augmentation policies that align the density of D_{train} with the density of the augmented $\mathcal{T}(D_{\text{valid}})$. This means that the transformations should either remove unimportant features or bolster meaningful features. In practice, FAA splits D_{train} into $D_{\mathcal{M}}$ and $D_{\mathcal{A}}$, where $D_{\mathcal{M}}$ is used to train the model and $D_{\mathcal{A}}$ is used to determine the policy via:

$$\mathcal{T}^* = \underset{\mathcal{T}}{\operatorname{argmin}} \mathcal{L}(\theta_{\mathcal{M}} | \mathcal{T}(D_{\mathcal{A}})) \quad (2)$$

where $\theta_{\mathcal{M}}$ is trained using $D_{\mathcal{M}}$. This approximates minimizing the distance between the density of $D_{\mathcal{M}}$ and $\mathcal{T}(D_{\mathcal{A}})$ by using augmentations to improve predictions with shared model parameters,

$\theta_{\mathcal{M}}$; see [25] for derivations. FAA obtains the final policy, \mathcal{T}^* , by exploring B candidate policies $\mathcal{B} = \{\mathcal{T}_1, \dots, \mathcal{T}_B\}$ with a Bayesian optimization method that samples a sequence of sub-policies from \mathcal{S} and adjusts the probabilities $\{p_1, \dots, p_{N_{\mathcal{T}}}\}$ and magnitudes $\{\lambda_1, \dots, \lambda_{N_{\mathcal{T}}}\}$ to minimize $\mathcal{L}(\theta|\cdot)$ on $\mathcal{T}(D_{\mathcal{A}})$ (see Appendix E for details). The top P policies from each data split are merged into \mathcal{T}^* . The network is then retrained using this policy on all training data, $\mathcal{T}^*(D_{\text{train}})$, to obtain the final network parameters θ^* . SelfAugment has three main differences that we discuss next.

Select the base policy: We need a *base augmentation policy* to perform the first pass of training to determine $\theta_{\mathcal{M}}$. We determine this policy by training for a short period of time on each of the individual transformations in \mathbb{O} as well as *random resize crop* (the top performing single transformation in [2]). We apply each transformation at every iteration, with $p = 1$ and a magnitude parameter λ stochastically set at each iteration to be within the ranges from [25] with uniform probability. We train each MoCo model until the loss curves separate—we found this to be around 10-15% of the total training epochs typically used for pretraining. We then freeze the MoCo network and train a linear image rotation predictor, ϕ_{rot} , for each network and evaluate it on held out training data. The base policy is then simply the transformation with the highest rotation prediction accuracy.

Searching augmentation policies: Given the base augmentations, we split the training data into k -folds. For each fold, we train a MoCo network, θ_{moco} , using the base augmentation. We then take the pretrained MoCo encoder, freeze the weights, and train a linear network, ϕ_{rot} to predict which rotation in $[0^\circ, 90^\circ, 180^\circ, 270^\circ]$ was applied to $\mathcal{D}_{\mathcal{M}}$. We use the same Bayesian optimization search strategy as FAA to determine the augmentation policies. However, as the loss function $\mathcal{L}(\theta|D)$ cannot use supervised accuracy as in FAA, we explore five variants of a self-supervised loss function and discuss their interpretation (Appendix F discusses and compares each of these loss functions):

- **Min. rotation error:** $\mathcal{T}^R = \text{argmin}_{\mathcal{T}} \mathcal{L}_{\text{rot}}(\theta_{\mathcal{M}}, \phi_{\text{rot}}|\mathcal{T}(D_{\mathcal{A}}))$ where \mathcal{L}_{rot} is the rotation cross-entropy loss, which yields policies that directly result in improved rotation prediction if the linear model was retrained on top of the same base network.
- **Min. InfoNCE:** $\mathcal{T}^{\text{I-min}} = \text{argmin}_{\mathcal{T}} \mathcal{L}_{\text{NCE}}(\theta_{\mathcal{M}}|\mathcal{T}(D_{\mathcal{A}}))$ where \mathcal{L}_{NCE} is the InfoNCE loss from Eq. 1, which yields policies that make it easier to distinguish image pairs in the contrastive feature space.
- **Max InfoNCE:** $\mathcal{T}^{\text{I-max}} = \text{argmin}_{\mathcal{T}} - \mathcal{L}_{\text{NCE}}(\theta_{\mathcal{M}}|\mathcal{T}(D_{\mathcal{A}}))$ negates the previous loss function, yielding policies that make it difficult to distinguish image pairs in the feature space, i.e. aligns the distribution of $\mathcal{D}_{\mathcal{A}}$ and $\mathcal{D}_{\mathcal{M}}$ in the feature space. We believe that this will result in a challenging contrastive loss task, that may be too challenging for the network to learn optimal representations.
- **Min \mathcal{L}_{rot} max \mathcal{L}_{NCE} :** $\mathcal{T}^{\text{minmax}} = \text{argmin}_{\mathcal{T}} (\mathcal{L}_{\text{rot}}(\theta_{\mathcal{M}}, \phi_{\text{rot}}|\mathcal{T}(D_{\mathcal{A}})) - \mathcal{L}_{\text{NCE}}(\theta_{\mathcal{M}}|\mathcal{T}(D_{\mathcal{A}})))$ yields policies with difficult transformations that maximize InfoNCE, while maintaining salient object features that minimize rotation error.

Retrain MoCo using the full training dataset and augmentation policy: We use the selected policy from each of the loss functions and retrain on the full dataset $\mathcal{D}_{\text{train}}$ from scratch. It is worth noting that because the augmentations we learn using SelfAugment are used for an instance contrastive task and not for rotation prediction or a correlated task such as supervised evaluation, the augmentations that minimize rotation loss are not necessarily the best augmentations for instance contrastive pretraining. Rather, this method provides the set of augmentations that would lead to high rotation prediction accuracy if we were to directly retrain the linear classifier on top of the original backbone. Hence, incorporating the InfoNCE loss balances the instance contrastive training task with the downstream task.

4 Experiments

Through the following experiments, we aim to establish that (i) image rotation prediction is highly correlated with the supervised performance of standard visual recognition tasks on common datasets and (ii) SelfAugment provides a competitive approach to augmentation selection, despite being fully unsupervised. All experiments were carried out using MoCo training with the standard Resnet-50 backbone [29] on 4 GPUs, using default training parameters from [3], see Appendix C. In Appendix D, we also conducted several auxiliary experiments exploring different architectures, self-supervised evaluation tasks, and fully connected layers for the rotation prediction.

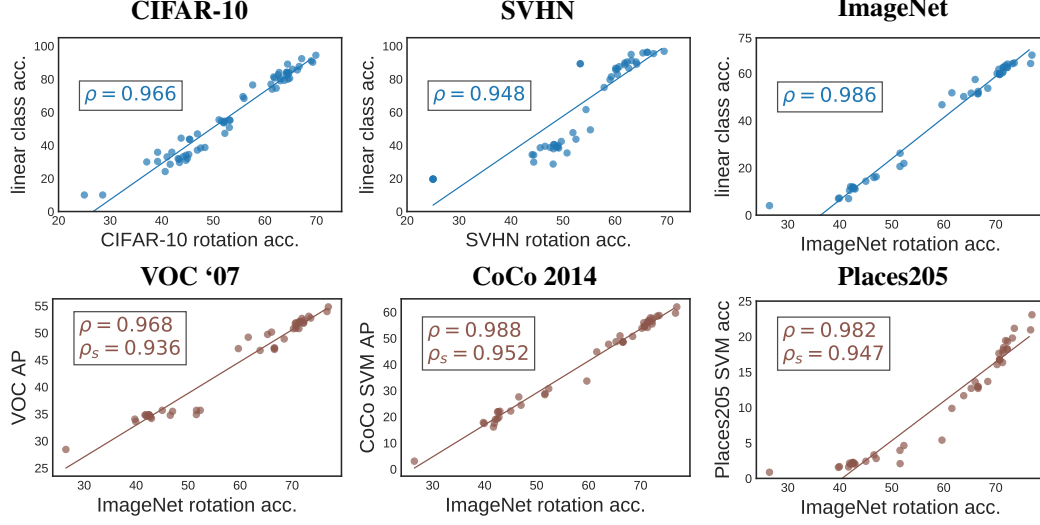


Figure 2: Top row: Correlation between supervised linear classification accuracy and linear image rotation prediction accuracy for three datasets. Bottom row: Correlation between rotation prediction and three transfer tasks from the ImageNet pretraining. ρ is the rank correlation between rotation prediction and the supervised task. ρ_s is the rank correlation between ImageNet supervised linear classification and the transfer task performance.

4.1 Supervised and self-supervised evaluation correlation

As evaluating all possible data augmentations for every dataset, training schedule, and downstream task is intractable, we evaluate a diverse sampling of RandAugment, SelfAugment, MoCoV2, and single-transform policies and training schedules for six benchmark datasets and downstream tasks:

CIFAR-10, SVHN, ImageNet [30, 31, 32]: For these benchmark datasets, we use augmentation policies: (i) random horizontal flip and random resize crop, (ii) RandAugment on top of (i) grid searched over parameters $\lambda = \{4, 5, 7, 9, 11\}$, $N_\tau = \{1, 2, 3\}$ at $\{100, 500\}$ epochs for CIFAR-10/SVHN and $\lambda = \{5, 7, 9, 11, 13\}$, $N_\tau = 2$ at $\{20, 60, 100\}$ epochs for ImageNet, (iii) each individual RandAugment transformations at 100 and 10 epochs for CIFAR-10 and ImageNet, respectively, (iv) scaling the magnitude λ from its min to max value for each N_τ at 500 epochs for CIFAR-10/SVHN and $\{20, 60, 100\}$ epochs for ImageNet, (v) for CIFAR-10/SVHN we also performed RandAugment using $K_T = \{3, 6, 9\}$ transformations at each of $\lambda = \{4, 7\}$ with $N_\tau = 2$, at 500 epochs, (vi) MoCoV2 augmentations at 100, 200 epochs for ImageNet, (vii) the five SelfAugment policies at 750 epochs for CIFAR-10/SVHN and 100 epochs for ImageNet. In total, this resulted in a diverse set of 61 models for CIFAR-10 and SVHN and 43 models for ImageNet.

Correlation results for linear classification: For CIFAR-10, SVHN, and ImageNet, we evaluate the frozen representations using the standard supervised linear classification and a self-supervised linear image rotation prediction classification; Appendix C contains all training details. The top row of Figure 2 shows the correlation between the supervised and self-supervised evaluations: CIFAR-10, SVHN, and ImageNet have respective Spearman rank correlations [33] of $\rho = (0.966, 0.948, 0.986)$ across the range of evaluated models—a very strong relationship between the standard supervised evaluation criteria and the proposed self-supervised evaluation criteria.

Figure 3 and Figure 4 demonstrate how rotation prediction can be used to select the augmentation transformations and their parameters. Figure 3 shows the performance of individual image transformation augmentations on SVHN. The left and middle plots show the classification accuracy compared with the InfoNCE loss function and top-1 contrastive accuracy (how well the self-supervised model predicts the augmented image pairs), while the right plot shows the rotation prediction accuracy for the fifteen transformations in ①, applied to SVHN and evaluated after 100 training epochs. Without the supervised feedback, selecting transformations with high or low values of InfoNCE or contrastive accuracy would miss the top performing transformation in the middle. By using the rotation predic-




Figure 3: For SVHN, we plot the supervised classification accuracy (y-axis) vs the InfoNCE loss function (left), contrastive top-1 accuracy (middle), and self-supervised linear rotation accuracy (right), for a self-supervised model trained using one of each transformation used by SelfAugment. Neither of the left two training metrics are a consistent measure of the quality of the representations, while the rotation prediction accuracy provides a strong linear relationship.

tion, each transformation has a clear linear relationship with the supervised performance, enabling the unsupervised selection of the best transformations.

Figure 4 demonstrates a similar result in that the individual image transformation parameters can be determined through rotation prediction. Specifically, we pretrained CIFAR-10 using the default MoCoV2 augmentation policy. Then, for each of the fifteen transformation in \mathbb{O} , we applied the transformation with probability 1.0 and a magnitude λ randomly selected between 0 and $\{0.25, 0.5, 0.75, 1.0\}$. We evaluated the individual transformation’s magnitude parameter using the supervised top-1 linear accuracy (classification) and self-supervised top-1 rotation prediction (rotation) as shown in Figure 4. Overall, the supervised linear classification and self-supervised rotation prediction select the the same magnitude parameter for 13 of the 15 transformations and have a strong Spearman rank correlation of 0.929. Taken together, these results indicate that rotation prediction can be used to both select individual transformations as well as the parameters of the transformations for an augmentation policy.

A central goal of representation learning is to learn transferable features. We therefore study the correlation between rotation prediction and the ImageNet transfer performance for the following:

- **VOC07 [34] object detection:** Following the specifics from [3], we transfer the pretrained models to a Faster R-CNN R50-C4 model and fine-tune all layers. Over three runs, we evaluate the mean of the default VOC metric, AP50, with COCO-style AP/AP₇₅ shown in Appendix D. Finetuning is performed on the `train2007+2012` set and evaluation is on the `test2007` set.
- **COCO2014 [35] image classification** Following [28], we train linear SVMs [36] on the frozen network and evaluate the accuracy over three end-to-end runs.
- **Places205 [37] low shot scene classification:** Following [28], we train linear SVMs on the frozen network using $k = \{1, 4, 8, 16, 32, 64\}$ labeled examples and evaluate the accuracy over five runs, with the average across all k used as the evaluation criteria, see Appendix D for a breakdown.

Correlation results for transfer performance: For VOC07, COCO2014, and Places205, the bottom row of Figure 2 shows the ImageNet rotation performance vs the transfer task performance, yielding strong rank correlations of $\rho = (0.968, 0.988, 0.982)$, respectively. For comparison, the rank correlation of the supervised linear classification on ImageNet is $\rho_s = \{0.936, 0.952, 0.947\}$. For each transfer task, the rotation correlation is stronger than the supervised correlation. In [4], the authors observe that “linear classification accuracy is not monotonically related to transfer performance in detection,” an observation that we find further evidence for across more transfer tasks. Further, we observe that *rotation prediction has a stronger transfer correlation than linear classification*.

4.2 Performance benchmarks

Here, we evaluate the SelfAugment and SelfRandAugment policies with the goal of establishing comparable results to the current state-of-the art methods without supervised feedback.

Setup: We evaluate: (i) linear classification performance on top of the frozen network using CIFAR-10, SVHN, and ImageNet, (ii) transfer performance of ImageNet models on PASCAL VOC object




Figure 4: Result of magnitude sweeps on CIFAR-10 for the 15 transformations optimized in Self-Augment. For each augmentation, we train MoCoV2 for 250 epochs using the usual MoCoV2 augmentations, and then added the single augmentation on top. We then vary the range of possible augmentation strength λ parameter to be between 0 and the value on the x-axis, and randomly select a value in that range, and apply the augmentation with 100 % probability. Rotation accuracy and classification accuracy have Spearman rank correlation $\rho = .929$, demonstrating how the relationship between rotation accuracy and classification accuracy can be used to select the parameters of individual transformations.

detection, COCO image classification, and Places205 few-show scene classification as described in the previous subsection, (iii) semi-supervised ImageNet top-1 accuracy with using only 1% or 10% of labels for linear training. All methods were evaluated with the same number of epochs and hyperparameters: 750 pretrain and 150 linear epochs for CIFAR-10/SVHN, 100 pretrain and 50 linear epochs for ImageNet, 24k finetuning iterations on VOC, and the exact training parameters/schedules from [28] for COCO and Places205. For SelfAugment, we used the hyperparameter settings from [25]: $P = 10$ policies, $T = 2$ transformations, and $K = 5$ folds of training. For SelfRandAugment, we used the grid search described in the previous subsection. See appendix C for all details.

Linear classification results: Table 1 compares all versions of SelfAugment to the current state-of-the-art performance from MoCoV2 policies [4] that were based on the extensive study of supervised policy evaluations in [2]. For linear classification with CIFAR-10, SVHN, and ImageNet, SelfAugment policies outperform MoCoV2’s policy. Where the largest gain occurs in the SVHN dataset, SVHN, consisting of images of house numbers, is the most distinct from ImageNet’s diverse, object centric images. Since MoCoV2’s policy is the result of extensive, supervised study on ImageNet, it

Table 1: Top-1 accuracy of SelfAug on CIFAR-10, SVHN, ImageNet, as well as ImageNet transfer tasks and semi-supervised experiments. An “R” superscript denotes best rotation accuracy. Bold results are greater than one standard deviation better across multiple runs, see Appendix C. Despite not using labels and not hand-tuning hyperparameters, SelfAug results in better performance than MoCov2 for 2 out of 3 benchmark datasets it was directly trained on, and comparable or improved transfer representations.

unsupervised	C10	SVHN	IN	VOC	COCO	Places	IN-1%	IN-10%
Base Aug	89.1	89.2	46.7	47.1	33.6	5.4	16.0	33.6
SelfRandAug	90.3	96.8^R	64.1	53.1	58.4	19.8	36.1	53.4
SelfAug (min rot)	91.0	94.9	57.4	50.2	50.9	13.6	26.4	45.1
SelfAug (min Info)	87.5	86.0	51.7	49.2	44.80	9.9	20.5	38.6
SelfAug (max Info)	90.1	96.2	63.3	52.6	57.8	19.4	34.4	52.7
SelfAug (minimax)	92.6^R	95.8	64.4	53.0	58.7	21.2	36.2	54.1
supervised feedback								
MoCoV2 [4]	92.3	96.4	64.2 ^R	54.0	59.6	20.8	37.9^R	54.9^R

is not surprising that it does not transfer as well to a distributionally distinct dataset such as SVHN. These results indicate that SelfAugment is a stronger approach to obtaining quality representations on distributionally distinct datasets, compared to directly using the optimized ImageNet policy.

Transfer and semi-supervised results: For the VOC07 object detection and COCO image classification transfer tasks, the MoCoV2 policy outperformed all other policies. The SelfAugment result had a stronger transfer performance for few-shot scene classification on the Places205 dataset across all k values (this difference exceeded one standard deviation at each k value; see Appendix C). Like ImageNet, VOC07 and COCO are natural images containing objects, while Places205 is a scene classification benchmark, consisting of complex scenes and diverse settings. SelfAugment’s policies, as with the linear classification, perform better for the dataset and task that substantially differs from ImageNet. While the SelfAugment policy outperformed MoCoV2 for the linear classification with 100% of the labels used for training, MoCoV2 performed better when using only 1% and 10% of the labeled data for training the linear classifier. These results indicate that using supervised evaluation to select a policy can lead to strong semi-supervised performance on the same dataset, but as indicated by the Places205 results, this policy may not transfer to other semi-supervised tasks.

Rotation prediction results: The mean rank correlation for the supervised linear classification and rotation prediction for all results in this subsection is $\rho = 0.978$, indicating that this correlation holds when removing the poorer performing models that were included in the previous subsection. For every linear classification result except ImageNet, the top performing augmentation policy also corresponded to the top performing rotation prediction performance, as indicated with an “R” superscript. For ImageNet, MoCoV2’s rotation prediction performance was significantly better than all other policies: $76.7 \pm 0.2\%$ compared to $73.6 \pm 0.2\%$ for SelfAugment’s minimax, over three linear evaluations. While for a similar analysis, the linear classification performance for MoCoV2 performed worse: $64.2 \pm 0.1\%$ compared to $64.4 \pm 0.1\%$. As shown in the previous subsection, however, the rotation prediction has a stronger correlation to transfer performance than the supervised linear classification, and indeed, that is the case here: on the VOC07 and COCO transfer tasks and the ImageNet 1% and 10% semi-supervised evaluations, the MoCoV2 policy significantly outperformed all other policies, while the SelfAug minimax policy had the best transfer performance for the Places205 task. In other words, across the transfer and semi-supervised evaluations, rotation prediction was more indicative of performance than supervised linear classification.

5 Discussion

If rotation prediction is highly correlated with supervised evaluations, why not directly train on it? As discussed in Appendix D.2, the authors of [21] found that rotation accuracy from training a full network on rotation prediction was only weakly correlated with supervised performance. Furthermore, as discussed in Appendix D.2, using a 2-layer MLP for rotation prediction drops the correlation from $\rho = .966$ to $.904$ even though rotation prediction accuracy improved. This indicates that while rotation prediction using a linear combination of the learned representations is an effective

evaluation, learning the representations by using rotation prediction does not strongly correlate with supervised performance.

Why doesn't minimizing rotation error or maximizing InfoNCE result in the best policies? We evaluate policies only *after* contrastive training, so minimizing rotation accuracy yields policies that should result in improved rotation prediction if we were to retrain the linear classifier. However, as indicated by the performance of the rotation error minimizing loss function, it is important to find a set of augmentations that also take the contrastive task into consideration. Similarly maximizing InfoNCE yields policies that make it difficult to distinguish image pairs in the feature space. In practice, this results in high magnitude augmentations that create a challenging contrastive task. However, the augmentations do not have any regularization that ensures the image maintains important features, leading to suboptimal performance. Taken together, maximizing InfoNCE produces challenging policies while minimizing rotation error provides regularization.

Which SelfAugment loss function should be used on a new, unlabeled dataset? *When training with a new unlabeled dataset, we recommend starting with the minimax SelfAugment loss function for augmentation policy selection.* This recommendation stems from its superior performance across all datasets and tasks. For comparison, the SelfAugment loss function minimizing the InfoNCE produced results that were often worse than the baseline policy, while both maximizing the InfoNCE or minimizing rotation prediction exceeded the baseline, but generally did not surpass MoCoV2's policy. In Appendix F, we show the effective magnitude of each individual transformation for each loss functions. We see that, as expected, the minimax loss function produces policies with intermediate magnitudes across all datasets. Its strong performance indicate that these intermediate magnitudes have found a *sweet spot* where the augmentations strike a balance between creating difficult instance contrastive tasks and retaining salient image features.

It is worth noting that across both SVHN and CIFAR-10, SelfAug with min InfoNCE loss function results in *worse* results than simply using the baseline augmentation. We believe this is because the augmentations selected in this case make it easier to tell samples apart using augmentations that do not actually add meaningful contrastive tasks - since they minimize InfoNCE - and are also not helpful towards improving the model's ability to generalize to new data. Accordingly, the model performs poorly when retrained with these augmentations. Similarly, it is somewhat surprising that we cannot just minimize rotation loss during SelfAug to achieve optimal results on the downstream task correlated with rotation - classification. We believe this is because the augmentations also need to take into account the contrastive loss. Simply ignoring the contrastive loss completely results in augmentations that may be too weak to provide a strong contrastive task. This is supported by the results in [2], where it is shown that using the augmentations from AutoAug [23], which are optimal for a supervised learning task, were less effective than a strong augmentation.

Computational efficiency: SelfAugment provides a highly efficient framework for finding augmentations for instance contrastive learning without using labels. As done in prior work [2], grid searching over various augmentation policies and values requires training many models to completion. Training a single instance of MoCoV2 on ImageNet using 4 V100 GPUs for 100 epochs takes 52 hours (208 GPU hours), which can be prohibitively expensive. SelfAugment finds useful, and typically improved, augmentation policies in just 730 GPU hours, and can learn additional augmentation policies in an additional 98 GPU Hours (see Appendix G for details). This is less than the computational cost of training just four full instances of MoCoV2. A single hyperparameter sweep for a single transformation can easily exceed four training iterations, meaning that Self-Augment provides a fast and efficient pathway to find a high performance set of augmentations for instance contrastive learning.

Future work: In this work, we have focused on establishing rotation prediction as a highly correlated task with downstream performance for self-supervised pretraining with natural images. In Appendix D, we perform a preliminary analysis of other self-supervised evaluation tasks, and in future work, further exploration of additional self-supervised tasks will likely yield promising results for other types of data (e.g. synthetic data and rotationally invariant data). By building upon this work, we hypothesize that exploring additional datasets and auto augmentation strategies will yield results that outperform hand-selected augmentation policies.

6 Conclusions

In this paper, we identified the problem that self-supervised representations are evaluated using labeled data, oftentimes using the labels from the “unlabeled” training dataset itself. We established that a self-supervised image rotation task is strongly correlated with the supervised performance for standard computer vision recognition tasks, and as a result, can be used to evaluate learned representations. Using this evaluation, we establish two unsupervised data augmentation policy selection algorithms and show that they can outperform or perform comparably to policies obtained using supervised feedback. We hope this work will encourage future research to leverage self-supervised evaluation criteria, as it opens up the wealth of supervised techniques that have been developed for making decisions when building and training models.

Acknowledgements Prof. Darrell’s group was supported in part by DoD, NSF as well as BAIR and BDD at Berkeley, and Prof. Keutzer’s group was supported in part by Alibaba, Amazon, Google, Facebook, Intel, and Samsung as well as BAIR and BDD at Berkeley. We would additionally like to thank Pieter Abbeel and Tete Xiao for helpful feedback and discussions. Wandb [38] graciously provided academic research accounts to track our experimental analyses.

References

- [1] Olivier J Hénaff, Aravind Srinivas, Jeffrey De Fauw, Ali Razavi, Carl Doersch, SM Eslami, and Aaron van den Oord. Data-efficient image recognition with contrastive predictive coding. *arXiv preprint arXiv:1905.09272*, 2019.
- [2] Ting Chen, Simon Kornblith, Mohammad Norouzi, and Geoffrey Hinton. A simple framework for contrastive learning of visual representations. *arXiv preprint arXiv:2002.05709*, 2020.
- [3] Kaiming He, Haoqi Fan, Yuxin Wu, Saining Xie, and Ross Girshick. Momentum contrast for unsupervised visual representation learning. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2020.
- [4] Xinlei Chen, Haoqi Fan, Ross Girshick, and Kaiming He. Improved baselines with momentum contrastive learning. *arXiv preprint arXiv:2003.04297*, 2020.
- [5] Zhirong Wu, Yuanjun Xiong, Stella X Yu, and Dahua Lin. Unsupervised feature learning via non-parametric instance discrimination. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 3733–3742, 2018.
- [6] H. Shin, M. Orton, D. J. Collins, S. Doran, and M. O. Leach. Autoencoder in time-series analysis for unsupervised tissues characterisation in a large unlabelled medical image dataset. In *2011 10th International Conference on Machine Learning and Applications and Workshops*, volume 1, pages 259–264, 2011.
- [7] X Zhang, Y Cui, Y Song, H Adam, and S Belongie. The imaterialist challenge 2017 dataset. In *FGVC Workshop at CVPR*, volume 2, page 3, 2017.
- [8] Jeff Donahue, Yangqing Jia, Oriol Vinyals, Judy Hoffman, Ning Zhang, Eric Tzeng, and Trevor Darrell. Decaf: A deep convolutional activation feature for generic visual recognition. In *International conference on machine learning*, pages 647–655, 2014.
- [9] Dumitru Erhan, Yoshua Bengio, Aaron Courville, Pierre-Antoine Manzagol, Pascal Vincent, and Samy Bengio. Why does unsupervised pre-training help deep learning? *Journal of Machine Learning Research*, 11(Feb):625–660, 2010.
- [10] Matthew D Zeiler and Rob Fergus. Visualizing and understanding convolutional networks. In *European conference on computer vision*, pages 818–833. Springer, 2014.
- [11] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep learning*. MIT press, 2016.
- [12] Yann LeCun, Yoshua Bengio, and Geoffrey Hinton. Deep learning. *nature*, 521(7553):436–444, 2015.
- [13] Alec Radford, Karthik Narasimhan, Tim Salimans, and Ilya Sutskever. Improving language understanding by generative pre-training.
- [14] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*, 2018.

- [15] Aaron van den Oord, Yazhe Li, and Oriol Vinyals. Representation learning with contrastive predictive coding. *arXiv preprint arXiv:1807.03748*, 2018.
- [16] Adam Coates and Andrew Y Ng. Learning feature representations with k-means. In *Neural networks: Tricks of the trade*, pages 561–580. Springer, 2012.
- [17] Spyros Gidaris, Praveer Singh, and Nikos Komodakis. Unsupervised representation learning by predicting image rotations. *CoRR*, abs/1803.07728, 2018.
- [18] Carl Doersch, Abhinav Gupta, and Alexei A Efros. Unsupervised visual representation learning by context prediction. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 1422–1430, 2015.
- [19] Mehdi Noroozi and Paolo Favaro. Unsupervised learning of visual representations by solving jigsaw puzzles. In *European Conference on Computer Vision*, pages 69–84. Springer, 2016.
- [20] Spyros Gidaris, Praveer Singh, and Nikos Komodakis. Unsupervised representation learning by predicting image rotations, 2018.
- [21] Alexander Kolesnikov, Xiaohua Zhai, and Lucas Beyer. Revisiting self-supervised visual representation learning. In *Proceedings of the IEEE conference on Computer Vision and Pattern Recognition*, pages 1920–1929, 2019.
- [22] Connor Shorten and Taghi M. Khoshgoftaar. A survey on Image Data Augmentation for Deep Learning. *Journal of Big Data*, 6(1):60, July 2019.
- [23] Ekin D Cubuk, Barret Zoph, Dandelion Mane, Vijay Vasudevan, and Quoc V Le. Autoaugment: Learning augmentation strategies from data. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 113–123, 2019.
- [24] Daniel Ho, Eric Liang, Ion Stoica, Pieter Abbeel, and Xi Chen. Population based augmentation: Efficient learning of augmentation policy schedules. In *ICML*, 2019.
- [25] Sungbin Lim, Ildoo Kim, Taesup Kim, Chiheon Kim, and Sungwoong Kim. Fast autoaugment. In *Advances in Neural Information Processing Systems*, pages 6662–6672, 2019.
- [26] Ekin D Cubuk, Barret Zoph, Jonathon Shlens, and Quoc V Le. Randaugment: Practical data augmentation with no separate search. *arXiv preprint arXiv:1909.13719*, 2019.
- [27] Chenxi Liu, Piotr Dollár, Kaiming He, Ross Girshick, Alan Yuille, and Saining Xie. Are labels necessary for neural architecture search? *arXiv preprint arXiv:2003.12056*, 2020.
- [28] Priya Goyal, Dhruv Mahajan, Abhinav Gupta, and Ishan Misra. Scaling and benchmarking self-supervised visual representation learning. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 6391–6400, 2019.
- [29] Christian Szegedy, Sergey Ioffe, Vincent Vanhoucke, and Alexander A Alemi. Inception-v4, inception-resnet and the impact of residual connections on learning. In *Thirty-first AAAI conference on artificial intelligence*, 2017.
- [30] Alex Krizhevsky, Vinod Nair, and Geoffrey Hinton. Cifar-10 (canadian institute for advanced research).
- [31] Ian J Goodfellow, Yaroslav Bulatov, Julian Ibarz, Sacha Arnoud, and Vinay Shet. Multi-digit number recognition from street view imagery using deep convolutional neural networks. *arXiv preprint arXiv:1312.6082*, 2013.
- [32] Olga Russakovsky, Jia Deng, Hao Su, Jonathan Krause, Sanjeev Satheesh, Sean Ma, Zhiheng Huang, Andrej Karpathy, Aditya Khosla, Michael Bernstein, Alexander C. Berg, and Li Fei-Fei. ImageNet Large Scale Visual Recognition Challenge. *International Journal of Computer Vision (IJCV)*, 115(3):211–252, 2015.
- [33] C Spearman. The proof and measurement of association between two things. *The American Journal of Psychology*, page 72, 1904.
- [34] M. Everingham, L. Van Gool, C. K. I. Williams, J. Winn, and A. Zisserman. The PASCAL Visual Object Classes Challenge 2007 (VOC2007) Results. <http://www.pascal-network.org/challenges/VOC/voc2007/workshop/index.html>.
- [35] Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C Lawrence Zitnick. Microsoft coco: Common objects in context. In *European conference on computer vision*, pages 740–755. Springer, 2014.

- [36] Bernhard E Boser, Isabelle M Guyon, and Vladimir N Vapnik. A training algorithm for optimal margin classifiers. In *Proceedings of the fifth annual workshop on Computational learning theory*, pages 144–152, 1992.
- [37] Bolei Zhou, Agata Lapedriza, Jianxiong Xiao, Antonio Torralba, and Aude Oliva. Learning deep features for scene recognition using places database. In *Advances in neural information processing systems*, pages 487–495, 2014.
- [38] Lukas Biewald. Experiment tracking with weights and biases, 2020. Software available from wandb.com.
- [39] Richard Zhang, Phillip Isola, and Alexei A Efros. Colorful image colorization. In *European conference on computer vision*, pages 649–666. Springer, 2016.

A Notation and definitions

Notation	Definition
ρ	Spearman rank correlation
θ_{moco}	MoCo encoder
ϕ_{rot}	Linear rotation predictor head
K	Number of folds used for training SelfAugment
T	Number of augmentations to apply to each image at each iteration
B	Number of policies to consider during Bayesian optimization search
P	Top number of policies to select from each fold in SelfAugment
\mathbb{O}	The set of candidate image transformations for an augmentation policy
\mathcal{O}	An image transformations in \mathbb{O}
\mathcal{S}	Set of sub-policies
τ	$\tau \in \mathcal{S}$ is the sequential application of N_τ consecutive transformation
N_τ	The number of consecutive transformations to apply in a sub-policy
λ	The magnitude of an image transformation
FAA	Fast AutoAugment, from [25]
InfoNCE	See Eq 1
MoCo	Momentum Contrast, from [3]

B Augmentation transformation details

Following [25, 26], we define the set of transformations used for SelfAugment and RandAugment, \mathbb{O} , as the PIL-based image transformations in Table 2. Each transformation has a minimum and maximum magnitude, λ , where for RandAugment, the entire range is discretized over 30 integers. See [26, 23] for further details and descriptions of each transformation.

C Additional training and experiment details

CIFAR-10, SVHN, ImageNet: Table 5 lists the training and experimental parameters for CIFAR-10, SVHN, and ImageNet. The training parameters were taken from [3] and adjusted for 4 GPUs, i.e. the learning rate and batch size were scaled by 0.5 since [3] experiments were conducted on 8 GPUs. Consult [3] and [4] for MoCo parameter information. For the linear classifier, we used 50 training iterations where the learning rate was $10\times$ reduced at 30 and 40 epochs for ImageNet and 20 and 30 epochs for CIFAR-10 and SVHN. In [3], the linear layer was trained over 150 iterations with a reduction at 80 and 100 iterations. In early experiments, we found the performance converged much earlier, and to reduce computational resources, we reduced all linear training iterations.

VOC07 Following [3], we transfer the ImageNet ResNet-50 weights to perform object detection using a Faster R-CNN R50-C4, with BN tuned. We fine-tuned all layers end-to-end. The image scale is [480, 800] pixels during training and 800 at inference. Training was on the VOC `trainval07+12` set and evaluation was on the `test2007` set. The R50-C4 backbone is similar to those available in Detectron2², where the backbone stops at the conv4 stage, and the box prediction head consists of the conv5 stage followed by a BN layer. Table 4 displays the AP/AP₅₀/AP₇₅ breakdown for three fine-tunings.

COCO2014/Places205: Following [28], we train Linear SVMs on frozen feature representations. We train a linear SVM per class for (80 for COCO, 205 for Places205) for the cost values $C \in 2^{[-19, -4]} \cup 10^{[-7, 2]}$. We used 3-fold cross-validation to select the cost parameter per class and then further calculate the mean average precision. The features are first normalized in a (N, 9k) matrix, where N

²<https://github.com/facebookresearch/detectron2>

Table 2: PIL image transformations used for SelfAugment and RandAugment. The min and max magnitude values are taken and from [23]

Name	Description	Min ($\lambda = 0.0$)	Max ($\lambda = 1.0$)
ShearX	shear the image along the horizontal axis with magnitude rate	-0.3	0.3
ShearY	shear the image along the vertical axis with magnitude rate	-0.3	0.3
TranslateX	translate the image in the horizontal direction by magnitude percentage	-0.45	0.45
TranslateY	translate the image in the vertical direction by magnitude percentage	-0.45	0.45
Rotate	rotate the image by magnitude degrees	-30	30
AutoContrast	adjust contrast so darkest pixel is black and lightest is white	0	1
Invert	invert the pixels of the image	0	1
Solarize	invert the pixels above a magnitude threshold	0	256
Posterize	reduce the number of bits for each color to magnitude	4	8
Contrast	adjust image contrast, where magnitude 0 is grey and magnitude 1 is original image	0.1	1.9
Color	adjust color of image such that magnitude 0 is black and white and magnitude 1 is original image	0.1	1.9
Brightness	brightness adjustment such that magnitude 0 is black image and 1 is original image	0.1	1.9
Sharpness	magnitude 0 is a blurred image and 1 is original image	0.1	1.9
Cutout	cutout a random square from the image with side length equal to the magnitude percentage of pixels	0	0.2
Equalize	equalize the image histogram	0	1

is number of samples in data and 9k is the resized feature dimension, to have norm=1 along each feature dimension. This normalization step is applied on evaluation data too. We use the following hyperparameter setting for training using LinearSVC sklearn class: `class_weight` ratio of 2:1 for positive:negative classes, `penalty=l2`, `loss=squared_hinge`, `tol=1e-4`, `dual=True` and `max_iter=2000`. Table 3 displays the Places205 results across five linear SVM trainings for all k values.

D Correlation study

In this section, we detail the full experimental setup, investigation, and results from our study of the correlation between rotation prediction performance with a linear network and supervised downstream task performance. We also include additional preliminary and ablation studies.

Table 3: Transfer results: Places205 top-1 accuracy for frozen Resnet50 encoder from ImageNet pretraining, with a linear SVM trained for scene classification, with $k = \{1, 4, 8, 16, 32, 64\}$ labeled training images for each class, averaged over five SVM trainings, where the errors indicate the standard deviation (see text for details).

Labeled samples	1	4	8	16	32	64
Base Aug	1.35 \pm .03	2.8 \pm 0.05	4.18 \pm .11	5.88 \pm .12	8.04 \pm .16	10.13 \pm .15
SelfRandAug	6.17 \pm .05	13.01 \pm .15	18.36 \pm .12	23.13 \pm .11	27.18 \pm .11	30.89 \pm .09
SelfAug (min rot)	3.58 \pm .06	7.85 \pm 0.09	11.77 \pm .10	15.60 \pm .17	19.66 \pm .10	23.26 \pm .05
SelfAug (min Info)	2.46 \pm .04	5.36 \pm 0.07	8.11 \pm .11	11.04 \pm .07	14.52 \pm .09	17.75 \pm .29
SelfAug (max Info)	5.87 \pm .06	12.73 \pm .14	17.88 \pm .16	22.64 \pm .21	26.95 \pm .12	30.56 \pm .17
SelfAug (minimax)	6.73 \pm .08	14.51 \pm .20	19.89 \pm .18	24.72 \pm .10	28.77 \pm .13	32.35 \pm .09
MoCoV2	6.65 \pm .11	14.06 \pm .13	19.59 \pm .14	24.57 \pm .13	28.56 \pm .08	32.24 \pm .10

Table 4: Transfer Result: For VOC07 test, this table reports the average AP50 and COCO-style AP/AP₇₅ over three runs of finetuning the Resnet50 encoder from ImageNet pretraining, where the errors indicate the standard deviation (see text for details).

Evaluation	AP	AP ₅₀	AP ₇₅
Base Aug	47.08 \pm 0.17	74.80 \pm 0.17	50.26 \pm 0.24
SelfRandAug	53.06 \pm 0.21	80.09 \pm 0.14	57.58 \pm 0.29
SelfAug (min rot)	50.13 \pm 0.20	77.66 \pm 0.18	53.94 \pm 0.21
SelfAug (min Info)	49.18 \pm 0.21	76.31 \pm 0.17	53.17 \pm 0.18
SelfAug (max Info)	52.66 \pm 0.18	79.69 \pm 0.15	57.33 \pm 0.25
SelfAug (minimax)	52.72 \pm 0.22	79.79 \pm 0.21	57.62 \pm 0.27
MoCoV2	53.94 \pm 0.23	80.64 \pm 0.18	59.34 \pm 0.31

D.1 Correlation study details

We study RandAugment, SelfAugment, and MoCoV2 augmentation policies, and then evaluate the performance using self-supervised rotation prediction with a linear layer. For CIFAR-10 and SVHN, we evaluate:

- A *base augmentation* of random left-right horizontal flips with $p = 0.5$ of being applied and random resize and crop transformation with magnitude range (0.2, 1.0), trained for 750 epochs.
- On top of the base augmentation³, we performed a RandAugment grid search with magnitude $\lambda = \{4, 5, 7, 9, 11\}$ and number of transformations applied $N_\tau = \{1, 2, 3\}$, evaluated at $\{100, 500\}$ epochs. We initially included $N_\tau = 4$ in the grid search, but this always led to a degenerate solution, whereby the training loss would not minimize the objective function and the evaluation would yield a chance result (25% self-supervised rotation prediction and 10% supervised classification for Cifar-10).
- Following [26], we also experimented with scaling the magnitude parameter λ from [4, 11] linearly throughout the training. We did this for each of the three N_τ values and evaluated the results at 500 epochs.
- As part of the SelfAugment algorithm, we trained an augmentation policy consisting of each of the fifteen individual transformations in RandAugment (transformations listed in §3.2). In addition, we include the random resize crop transformation, as it was shown to be the best performing individual transformation in [2]. The magnitude for each transformation was stochastically selected from its magnitude range defined in §B at each iteration. Each of these 16 transformations were evaluated after a short training cycle of 100 epochs. Each of these transformations were applied on top of a random left-right horizontal flip applied with $p = 0.5$.
- Using the rotation prediction results from each of the individual transformation policies, we selected the top $K_T = \{3, 6, 9\}$ augmentations, and trained RandAugment using only these

³Using this base augmentation systematically improved all RandAugment results over not using a base augmentation, both in terms of its rotation prediction performance and supervised linear classification performance.

Table 5: Detailed training parameters for CIFAR-10, SVHN, and ImageNet experiments carried out in this paper.

MoCo Params	CIFAR 10/SVHN	ImageNet
Batch Size	512	128
moco-dim	128	128
moco-k	65536	65536
moco-m	0.999	0.99
moco-t	0.2	0.2
num-gpus	4	4
lr	0.4	0.015
schedule	120, 160	120, 160
momentum	0.9	0.9
weight decay	1e-4	1e-4
Linear Classifier Params		
lr	15	30
batch size	256	256
momentum	0.9	0.9
weight decay	0.0	0.0
schedule	20, 30	30, 40
epochs	50	50
Rotation Classifier Params		
lr	15	30
batch size	256	256
momentum	0.9	0.9
weight decay	0.0	0.0
schedule	20, 30	10, 20
epochs	50	50

transformations. We performed this training for $\lambda = \{4, 7\}$ with $N_\tau = 2$, evaluated at 500 epochs.

- We further included the five SelfAugment policies from each of its loss functions, evaluated at 750 epochs.

In total, this resulted in 61 different models for each of CIFAR-10 and SVHN. For ImageNet we evaluated:

- A *base augmentation* of random left-right horizontal flips with $p = 0.5$ of being applied and random resize and crop transformation with magnitude range $(0.2, 1.0)$, trained for 100 epochs.
- On top of the base augmentation⁴, we performed a RandAugment grid search with magnitude $\lambda = \{5, 7, 9, 11, 13\}$ and number of transformations applied $N_\tau = 2$, evaluated at $\{20, 60, 100\}$ epochs.
- Following [26], we also experimented with scaling the magnitude parameter λ from $[5, 13]$ linearly throughout the training. We did this for $N_\tau = 2$ and evaluated the results at 100 epochs.
- As part of the SelfAugment algorithm, we trained an augmentation policy consisting of each of the fifteen individual transformations in RandAugment (transformations listed in §). In addition, we include the random resize crop transformation, as it was shown to be the best performing individual transformation in [2]. The magnitude for each transformation was stochastically selected from its magnitude range defined in §B at each iteration. Each of these 16 transformations were evaluated after a short training cycle of 100 epochs. Each of

⁴Using this base augmentation systematically improved all RandAugment results over not using a base augmentation, both in terms of its rotation prediction performance and supervised linear classification performance.

these transformations were applied on top of a random left-right horizontal flip applied with $p = 0.5$.

- We further included the five SelfAugment policies from each of its loss functions, evaluated at 100 epochs.
- To further compare with state-of-the-art models, we compare with the MoCoV2 augmentation policy evaluated at 100, 200 epochs

In total, this resulted in 43 different models for ImageNet.

Supplementing the main correlation results shown in the experiments section, Figure 5 shows the the ImageNet rotation prediction correlations with transfer performance to VOC07 for all AP, AP₅₀, and AP₇₅ evaluations. Figure 6 the ImageNet rotation prediction correlations with transfer performance to Places205 few label scene classification using $k = \{1, 4, 8, 16, 32, 64\}$ labels per scene class. For both VOC07 and Places205, the Spearman rank correlation with rotation prediction is indicated with ρ , while the rank correlation with the supervised ImageNet classification performance is indicated with ρ_s . Across all evaluation metrics, the rotation correlation is higher than the supervised correlation.




Figure 5: The ImageNet rotation prediction correlations with transfer performance to VOC07 for all AP, AP₅₀, and AP₇₅ evaluations. The Spearman rank correlation with rotation prediction is indicated with ρ , while the rank correlation with the supervised ImageNet classification performance is indicated with ρ_s .

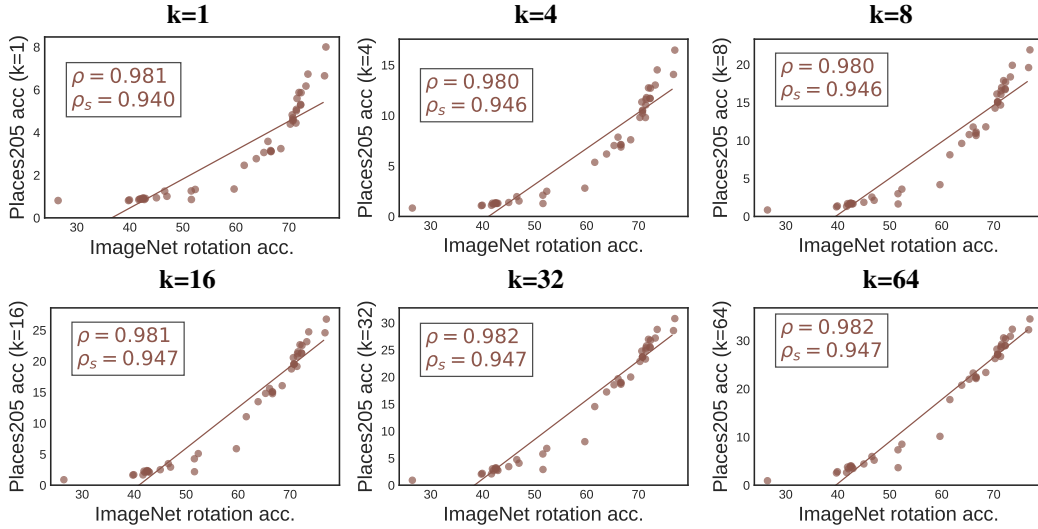


Figure 6: The ImageNet rotation prediction correlations with transfer performance to Places205 few label scene classification using $k = \{1, 4, 8, 16, 32, 64\}$ labels per scene class. The Spearman rank correlation with rotation prediction is indicated with ρ , while the rank correlation with the supervised ImageNet classification performance is indicated with ρ_s .

D.2 Correlation with an MLP rotation prediction

When following the same evaluation protocol, except using a 2 layer MLP instead of a linear layer for rotation prediction, we find that the CIFAR-10 correlation with the supervised linear classification

drops from 0.966 to 0.904. This correlation drop indicates that using a simple linear layer, rather than a more complicated network, is ideal for evaluation of the learned representations. A more complicated network can learn its own representation, which distances the evaluation from the learned representations.

D.3 Correlation across architectures

In [21], the authors showed that when training the entire network to recognize image rotations, rather than just a linear layer, the rotation prediction performance did not correlate across architectures. To examine this question for a linear rotation prediction layer, we study the CIFAR-10 correlation for three encoding architectures: ResNet18, ResNet50, Wide-ResNet-50-2, using RandAugment with a grid search over the number of applied transformations $N_\tau = \{1, 2, 3\}$ and magnitudes $\lambda = \{4, 5, 7, 9, 11\}$ evaluated at 100 epochs. Figure 7 shows the results for each architecture: the overall Spearman rank correlation across all architectures is $\rho = 0.921$, which is in between the Wide-ResNet-50-2 and ResNet18 Spearman correlations of 0.924 and 0.914 and less than the ResNet50 correlation of 0.957.

Overall, these strong correlations indicate that the rank correlation for a linear rotation prediction is maintained across architectures. Given the drop in performance when using a 2 layer MLP for rotation prediction instead of a linear layer, combined with the lack of correlation discovered when using a full network in [21], we conclude that using a linear layer to predict rotations is important to disentangle the learned representations from the ability of the network to learn its own rotation features.




Figure 7: The Spearman rank correlation for CIFAR-10 RandAugment grid search for ResNet18, ResNet50, and Wide-ResNet-50-2, as well the combined rank correlation.

D.4 Correlation for other self-supervised tasks

In addition to the rotation prediction study presented in this paper, we also performed a preliminary study for self-supervised evaluation methods of solving jigsaw puzzles [19] and predicting the colorization of the image [39]. For solving jigsaw puzzles, the input image is divided into four quadrants and randomly shuffled to one of 24 permutations; the pretext task is then to predict the original permutation given the shuffled version. For image colorization, the input image is converted into a greyscale image and the prediction task is then a pixel-wise regression of each pixel. Following rotation prediction, we train a linear layer for both jigsaw prediction and colorization on top of the frozen encoder network.

As a preliminary analysis, we randomly select 30 augmentation policies from the set of policies described in Appendix 2 for CIFAR-10 and SVHN. We then evaluate the self-supervised task and a top-1 supervised linear classification as described in §3.2. We use the Spearman rank correlation between the self-supervised and supervised tasks as a measure of the effectiveness of the self-supervised evaluation. Table 6 shows the Spearman rank correlation with the supervised linear classification for all three self-supervised evaluation tasks on both CIFAR-10 and SVHN. For both CIFAR-10 and SVHN, rotation prediction with a linear layer has a significantly stronger correlation with the

supervised linear classification compared to the jigsaw and colorization tasks. This preliminary experiment directed the attention of our study to rotation prediction, though we note that various alterations to the self-supervised task formulation (such as using a MLP instead of a linear network or evaluating on synthetic images) may show that a different self-supervised evaluation task is most highly correlated with supervised performance. We leave such investigation to future work.

Table 6: As a preliminary analysis, we randomly select 30 augmentation policies from the set of policies described in Appendix 2 for CIFAR-10 and SVHN. We then evaluate the specified self-supervised tasks (rotation, jigsaw, colorization) and a top-1 supervised linear classification as described in §4. We use the Spearman rank correlation (ρ) between the self-supervised and supervised tasks as a measure of the effectiveness of the self-supervised evaluation.

Self-supervised evaluation	CIFAR-10 (ρ)	SVHN (ρ)
Rotation	0.969	0.941
Jigsaw	0.926	0.902
Colorization	0.639	0.814

D.5 Correlation example: tuning transformation parameters

Using CIFAR-10 and the MoCoV2 augmentation policy, we freeze the augmentation parameters except for the minimum resize value of RandomResizeCrop, which we set as $\{0, 0.01, 0.08, 0.20, 0.50, 0.80\}$. Figure 8 shows the supervised classification evaluation, self-supervised rotation evaluation, and the correlation between the two evaluations. Using a linear rotation prediction evaluation finds the same optimal value as the supervised evaluation, 0.20.




Figure 8: For the RandomResizeCrop transformation with CIFAR-10, we set the minimum resize value to be $\{0, 0.01, 0.08, 0.20, 0.50, 0.80\}$, and show the supervised linear classification performance on the left plot, self-supervised linear rotation prediction performance on the middle plot, and relationship between the supervised and self-supervised evaluation on the right plot. Note that the linear rotation prediction performance perfectly matches the supervised performance and can be used to select the parameter value.

E Augmentation policy exploration via Bayesian optimization

As in [25], we used policy exploration search. Since there are an infinite number of possible policies, we applied Bayesian optimization to explore augmentation strategies. In line 7 in Algorithm 1, we employed the expected improvement (EI) criterion acquisition function to explore \mathcal{B} efficiently: $EI(\mathcal{T}) = \mathbb{E}[\min(\mathcal{L}(\theta, \phi | \mathcal{T}(\mathcal{D}_{\mathcal{A}}) - \mathcal{L}^{\dagger}, 0)]$. Here \mathcal{L}^{\dagger} represents a constant threshold determined by the quantile of observations amongst previously explored policies. As in [25], we used variable kernel density estimation on a graph-structured search space to approximate the criterion. Since this method is already implemented in the tree-structured Parzen estimator algorithm we used Ray⁵ and Hyperopt to implement this in parallel.

The full list of augmentations and range of magnitudes explored during augmentation policy exploration are detailed in Table 2.

⁵<https://github.com/ray-project/ray>

F SelfAugment Loss Functions

In this section we discuss the loss functions used for SelfAug in greater detail and their resulting augmentation policies which are summarized in Figure 9.

- **Min. rotation error:**

$$\mathcal{T}^R = \operatorname{argmin}_{\mathcal{T}} \mathcal{L}_{\text{rot}}(\theta_{\mathcal{M}}, \phi_{\text{rot}} | \mathcal{T}(D_{\mathcal{A}}))$$

where \mathcal{L}_{rot} is the rotation cross-entropy loss, which yields policies that should result in improved rotation prediction if we were to retrain the linear classifier with the augmentations provided. However, since the augmentations are instead used to create a contrastive learning task, it is important that we find a different set of augmentations that take the contrastive task into consideration. Indeed, using this loss function results in weak expected augmentation strengths (see Figure 9), resulting in relatively poor downstream performance (see Table 1).

- **Min. InfoNCE:**

$$\mathcal{T}^{\text{I-min}} = \operatorname{argmin}_{\mathcal{T}} \mathcal{L}_{\text{NCE}}(\theta_{\mathcal{M}} | \mathcal{T}(D_{\mathcal{A}}))$$

where \mathcal{L}_{NCE} is the InfoNCE loss from Eq. 1, which yields policies that make it easier to distinguish image pairs in the contrastive feature space. This loss function should result in small magnitude augmentations, since a trivial way to minimize InfoNCE is to apply no transforms - resulting in trivial minimization of the InfoNCE loss. Indeed, the loss function yields (i) lower expected augmentation strengths and (ii) emphasizes transformations like *Sharpness*, Figure 9. We did not expect this loss function to perform well, and was primarily included as a sanity check that minimizing InfoNCE would result in light augmentations.

- **Max InfoNCE:**

$$\mathcal{T}^{\text{I-max}} = \operatorname{argmin}_{\mathcal{T}} - \mathcal{L}_{\text{NCE}}(\theta_{\mathcal{M}} | \mathcal{T}(D_{\mathcal{A}}))$$

negates the previous loss function, yielding policies that make it difficult to distinguish image pairs in the feature space. In practice, this results in high magnitude augmentations and emphasizes augmentations like *Invert*. We hypothesized that maximizing InfoNCE could result in strong augmentations that could create a challenging contrastive task. However, the augmentations do not have any regularizing that would ensure the image maintains important features, leading to relatively suboptimal performance (see Table 1).

- **Min \mathcal{L}_{rot} max \mathcal{L}_{NCE} :**

$$\mathcal{T}^{\text{minmax}} = \operatorname{argmin}_{\mathcal{T}} (\mathcal{L}_{\text{rot}}(\theta_{\mathcal{M}}, \phi_{\text{rot}} | \mathcal{T}(D_{\mathcal{A}})) - \mathcal{L}_{\text{NCE}}(\theta_{\mathcal{M}} | \mathcal{T}(D_{\mathcal{A}})))$$

yields policies with difficult transformations that maximize InfoNCE, while maintaining salient object features that minimize rotation error. These augmentations were found to be the most useful for contrastive learning. In practice, we normalized \mathcal{L}_{rot} and \mathcal{L}_{NCE} by their expected value for the training data across the K-folds when training with the base augmentation, since \mathcal{L}_{NCE} was usually higher than \mathcal{L}_{rot} , but had each loss contribute equally for augmentation policy selection. Policies that optimized this loss emphasized transformations like *Equalize*, *AutoContrast* and *Contrast* more than others. For ImageNet, these transformations proved to be challenging yet useful. Notably, they closely resemble the effects of MoCov2 and SimCLR’s Color Jitter [4].

Finally, while **min \mathcal{L}_{rot} max \mathcal{L}_{NCE}** works well, one could potentially gain performance by weighting the importance of minimizing \mathcal{L}_{rot} vs maximizing \mathcal{L}_{NCE} . Hence we propose experimenting with different values of λ_{NCE} and λ_{rot} when optimizing the following objective: **min $\lambda_{\text{rot}} \mathcal{L}_{\text{rot}}$ max $\lambda_{\text{NCE}} \mathcal{L}_{\text{NCE}}$** .

G Computational Efficiency

SelfAug can be broken up into three steps:

1. **Finding the base augmentation.** This entails training 16 instances of MoCoV2 for 10-15% of the total epochs typically used for pretraining, using all of the training data available.
2. **Training on K-Folds using the base augmentation.** For CIFAR-10/SVHN we used the entire training dataset for this step, and evaluated for the same number of epochs as we use to train the final models, using all the training data from each fold. For ImageNet, we used a subset of 50,000 images and trained for 5x longer, to make up for the smaller amount of images. This was done to improve the computational efficiency of step 3.




Figure 9: Visualization of the augmentation policies found with SelfAugment as well as Fast AutoAugment for supervised learning. This figure shows expected augmentation strength (mean magnitude \times normalized probability) of each augmentation, evaluated across all datasets and loss functions. As expected, minimizing InfoNCE results in augmentations with smaller magnitude, and emphasizes augmentations that do not alter the image much (e.g. *Sharpness*). Maximizing InfoNCE results in augmentations with a larger magnitude and emphasizes augmentations (e.g. *Invert*, *Solarize*) which heavily alter the image. The minimax loss function yields an augmentation policy that strikes a middle ground between the two, with strong augmentations, but reduced emphasis on heavy augmentations like *Invert*. Comparison of the policies that worked best with contrastive learning relative to the original FAA policies reveals that our policies are (i) stronger, and (ii) have more variability in a single augmentation’s $E(\lambda * p)$ relative to policies that were used for supervised learning.

3. **Finding augmentation policies.** Because we only need to complete forward passes of the network, this is relatively efficient. We use the held out data from each of the K-folds to evaluate the model with different augmentation policies applied.

It is important to note that steps 1 and 2 are *shared* across all augmentation policies found with SelfAug, meaning that they do not need to be repeated to find a new augmentation policy using a different loss function. This could allow for rapid experimentation with new loss functions for SelfAug in the future.

Meanwhile, SelfRandAug’s computational time is completely determined by the user-defined search space over λ and N_τ . However, because it requires training multiple models to completion on the entire dataset, it is more computationally intensive than SelfAug. SelfRandAug’s computation time could be reduced by training on a subset of Images in a large training set.

Finally, it is worth noting that evaluating a *single* augmentation’s various hyperparameters for MoCoV2 can be extremely computationally expensive. Because pretraining MoCoV2 for 100 epochs then training a linear head for evaluation takes a total of 244 GPU Hours on a Tesla V100 GPU, searching over various augmentations and their strength and probability parameters can easily require thousands of GPU hours. Hence, an additional contribution of this work is providing a fast, automatic method for selecting augmentations for instance contrastive learning.

Table 7: Computational time, measured in one hour on one NVIDIA Tesla V100 GPU for Self-Augment and SelfRandAug on ImageNet. SelfAug times are evaluated when using the minimax loss function, which takes the most time because we evaluate both InfoNCE and rotation loss. For SelfRandAug, computational time reflects our grid search over $\lambda = \{5, 7, 9, 11, 13\}$, $N_\tau = 2$, plus the time to train a rotation head on top of each representation. This time could be larger or smaller depending on the parameters λ and N_τ a user wants to search over.

Algorithm	Find Base Augmentation	Train K-Folds	Find Augmentations	Total
SelfAug	476.6	155.0	97.7	729.4
SelfRandAug				1220.0